

Multi-Agent A*

AAAI 2015 Submission 2327

Abstract

abstract

Notes

PAD* notes: should we check last element of CLOSED instead of first in OPEN?

For $i = 1, \dots, N$, agent i has a graph defined by vertex set V_i and edge weights $c_i : V_i \times V_i \rightarrow \mathbb{R}^+$. We consider constraints Φ , each defining a relation on the relative visit times of two sets of nodes $\Phi^-, \Phi^+ \subset \cup_i V_i$. There are four constraint types:

- $\min \leq \min$ (opening a door)
- $\max \leq \min$ (closing a door)
- $\max \leq \max$ (restoration)
- $\min \leq \max$ (sequence)

Given individual plans for each agent, an optimal multi-agent plan can be constructed greedily forward in time, when all constraints are of the opening or closing kind. An optimal multi-agent plan can be constructed greedily backward in time when all constraints are of the closing or restoration kind. Since restoration and sequence constraints can be represented in terms of opens and closes, we consider the former as the only types without loss of generality.

Representing restoration constraints in terms of opens and closes: suppose B restores (i.e. cleans after) A. Then create a copy B' of B. Let B' close A, and B' open any agent's goal node. To deal with the boundary condition where neither A nor B are visited, the start node should have a similar copy, reachable at zero cost.

Representing sequence constraints in terms of opens: suppose (A, B) are a mandated sequence. Then create a copy B' of B. Let A open B', and B' open any agent's goal node.

NP-completeness proof:

Let's say there are m clauses and the variables are x_1, x_2, \dots, x_n . Draw the edges $(start_1, x_1)$, $(start_1, \neg x_1)$, (x_i, x_{i+1}) , $(x_i, \neg x_{i+1})$, $(\neg x_i, x_{i+1})$, $(\neg x_i, \neg x_{i+1})$, $(x_n, goal_1)$, $(\neg x_n, \neg goal_1)$. For the other agent, draw the edges

$(start_2, c_0)$, $(c_{i-1}, req_{i,j})$, $(req_{i,j}, c_i)$, $(c_m, goal_2)$. Here, $req_{i,j}$ is a door that needs to be opened by the node corresponding to the j 'th variable of the i 'th clause.

$g_{s_1, \dots, s_N}(s_i)$ is the minimum time by which agent i can reach state s_i , assuming every agent j will go to s_j along some known path (following $bp(s_j)$, say) and wait there forever. Note that the state space for each agent is the Cartesian product of its graph's vertex set, and the set of partial permutations of constraints.

Define $f(s_1, \dots, s_N) = \max_i g_{s_1, \dots, s_N}(s_i) + \epsilon h(s_i)$. This key is impractical to compute over all possible tuples. Further approximations will need to be made. Note that (s_1, \dots, s_N) range over the goal nodes as well as the OPEN list.

Even if we could compute f , does it guarantee $N\epsilon$ optimality? To prove that, it would suffice to show that we only expand nodes whose individual-agent g -values are ϵ -optimal. However, the logic fails because passing through a door-opening node can result in a sudden drastic decrease in the g_{s_1, \dots, s_N} value, due to constraints being lifted. To remedy this, instead of defining g by "wait-forever" semantics, we could treat doors as being opened after the partial paths to (s_1, \dots, s_N) , at some lower bound time.

A lower-bound for $f(s_1, \dots, s_n)$ could be computed as $f(s_j)$ by building a coarse search tree at the level of constraint sequences, where a macro-expansion consists of searching toward each possible next constraint node.

Want to prove that $g(s)$ is ϵ -optimal at expansion (when its jointPriority value is minimal), and thus the joint paths are $N\epsilon$ optimal. Suppose $g(s) > \epsilon g^*(s)$. As induction hypothesis, let $s' \in OPEN$ be such that $g(s') \leq \epsilon g^*(s')$. Fix a priority-minimizing tuple $(s_1, \dots, s', \dots, s_n)$ for s . Now consider the tuple $(s_1, \dots, s', \dots, s_n)$. If its priority were less, it would follow that s' will be expanded before s . But, path to s might open doors much sooner than path to s' !

Pending further insights, it seems the state is forced to remember sequences of constrained nodes instead of merely sequences of constraints. This way, we guarantee prefix-suboptimality, and as a bonus get global ϵ -suboptimality without the extra factor of N .

Algorithm 1 *MultiAgentA*HighLevel()*

```

 $\forall i : s_i := start_i$ 
while  $g(goal_1, \dots, goal_N) > f(s_1, \dots, s_N)$  do
   $\forall i : expand(s_i)$ 
   $(s_1, \dots, s_N) = \arg \min_{s_1, \dots, s_N} f(s_1, \dots, s_N)$ 
end while
```

Algorithm 2 *MultiAgentA*()*

```

 $\forall i : s_i := start_i$ 
 $OPEN := \{start_1, \dots, start_N\}$ 
while  $jointPriority(goal_1, \dots, goal_N) >$ 
 $OPEN.fmin()$  do
   $expand(OPEN.extractMin())$ 
end while
```

Algorithm 3 *expand(s)*

```

for all  $s' \in successors(s)$  do
  if  $g(s') > g(s) + c(s.node, s'.node)$  then
     $g(s') := g(s) + c(s.node, s'.node)$ 
     $bp(s') := s$ 
    insert  $s'$  in  $OPEN$ 
  end if
end for
```

Algorithm 4 *f(s)*

```

 $f := \infty$ 
for all tuples  $s_{1\dots N} \in OPEN$  that include  $s$  do
   $f := \min(f, jointPriority(s_{1\dots N}))$ 
end for
return  $f$ 
```

Algorithm 5 *successors(u, seq_u)*

```

succlist :=  $\emptyset$ 
for all  $v$  such that  $c(u, v) < \infty$  do
   $seq_v := seq_u$ 
  for all min constraints  $\Phi^\pm \notin seq_u$  with  $v \in \Phi^\pm$  do
    append  $\Phi^\pm$  to  $seq_v$ 
  end for
  for all max constraints  $\Phi^\pm$  with  $v \in \Phi^\pm$  do
    remove  $\Phi^\pm$  from  $seq_v$ 
    append  $\Phi^\pm$  to  $seq_v$ 
  end for
  push  $(v, seq_v)$  onto succlist
end for
return succlist
```

Algorithm 6 *greedyCost(π_1, \dots, π_n)*

```

 $\forall \Phi : trigger(\Phi) := \Phi$  is open or  $\Phi^-$  is absent in  $\pi_{1\dots N}$ 
 $\forall i \in 1 \dots N : u_i := v_i := start_i$ 
 $events := \{0 : \{1 \dots N\}\}$ 
while  $events$  is not empty do
   $(time, ids) := extractMin(events)$ 
  for all  $i \in ids$  do
     $u_i := v_i$ 
    for all constraints  $\Phi^-$  with  $v \in \Phi^-$  do
      if  $\Phi$  is an open constraint or  $\Phi^-$  appears in no
      suffix of  $\pi_{1\dots N}$  after  $u_{1\dots N}$  then
         $trigger(\Phi) := true$ 
      end if
    end for
  end for
  for all  $i \in 1 \dots N$  with  $u_i = v_i \neq goal_i$  do
     $v_i :=$  successor of  $u_i$  on  $\pi_i$ 
    for all constraints  $\Phi^+$  with  $v \in \Phi^+$  do
      if  $\neg trigger(\Phi)$  then
         $v_i := u_i$  and continue with the next  $i$ 
      end if
    end for
     $events(time + c(u_i, v_i)).insert(i)$ 
  end for
end while
if  $u_{1\dots N} = goal_{1\dots N}$  then
  return  $time$ 
else
  return  $\infty$ 
end if
```

Algorithm 7 *jointPriority(s_1, \dots, s_N)*

```

build  $\pi_{1\dots N}$  by following  $bp()$  from  $s_{1\dots N}$ 
for all  $i \in 1 \dots N$  such that  $s_i.node \neq goal_i$  do
  extend  $\pi_i$  with a zero cost edge to a node that opens
  all doors the agent can open, followed by  $goal_i$  at cost
   $\epsilon h(s_i)$ 
end for
return  $greedyCost(\pi_1, \dots, \pi_n)$ 
```
